


Attack Success Rate (LLM-as-judge) FinanceBench 58-query Clean-Correct Subset



Benign accuracy	100% (58/58 clean-correct subset)
Poison retrieval	100% (290/290 condition runs)
Source corpus	FinanceBench annual reports and filings
Evaluation set	58 queries, 58 optimized attack artifacts
Index / retriever	BAAI/bge-small-en-v1.5, top-k = 5
Architectures	single-agent, orchestrator (3), debate (3)
Runs	targeted / global, 290 rows gpt-4o-mini
ASR definition	harmful phrase match OR LLM target-claim judge