

Experiment Environment — Cybersecurity Corpus

Parameter	Value
Source documents	4 PDFs (NIST CSF 2.0, SP 800-53 Rev5, SP 800-61 Rev3, CISA IR Playbook)
Index chunks	1,731
Chunk size / overlap	384 tokens / 64 tokens
Embedding model	BAAI/bge-small-en-v1.5 (local)
Top-k per agent	5
Queries	8 (q001–q008)
Attack artifacts	8 (one per query, HotFlip/AgentPoison)
Trigger length	5 adversarial tokens
Target claims	False NIST/CISA policy statements
Systems evaluated	Single-agent · Orchestrator (3 agents) · Debate (3 agents)
Conditions	Clean · Targeted (1 poisoned agent) · Global (all agents)
Trials per condition	3
Generation model	gpt-4o-mini
Orchestrator mode	choose_best (LLM pick)
Debate mode	majority_vote · max 4 rounds · stable_for=2
Total runs	339 (clean: 173 · attack: 166)
ASR scorer	LLM-as-judge (GPT-4o-mini)